



DD2437 – Artificial Neural Networks and Deep Architectures (annda)

Lecture 1: **Course introduction and fundamental concepts**

Pawel Herman and Erik Fransén

Computational Science and Technology (CST)

KTH Royal Institute of Technology



Overview

- A historical note
- Course outline
 - Learning activities
 - Assessment, grading
- Introduction to ANNs
 - Fundamental characteristics
 - Brain inspirations
 - Current developments – theory and applications

- **A historical note**
- Course outline
- Introduction to ANNs

Brief history of ANNs


- McCulloch and Pitts
- Donald Hebb “The Organization of Behaviour”



- A historical note
- Course outline
- Introduction to ANNs

Brief history of ANNs


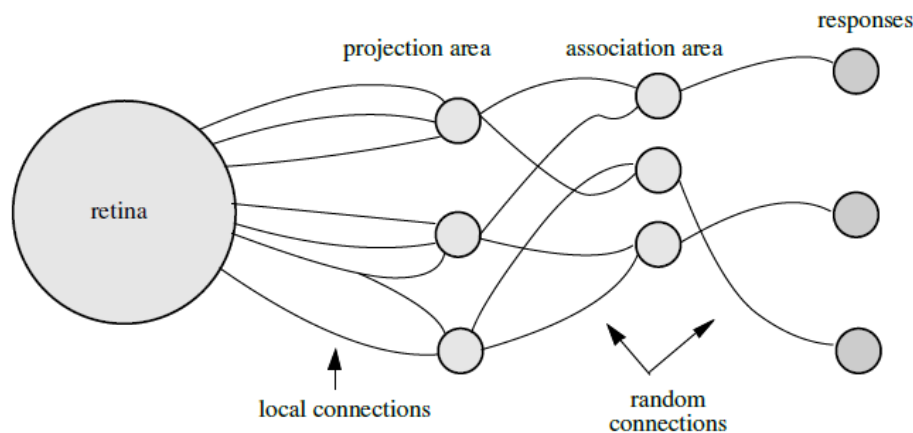


- McCulloch and Pitts
- Donald Hebb “The Organization of Behaviour”
- ADALINE for binary patterns by Widrow and Hoff



- A historical note
- Course outline
- Introduction to ANNs

Brief history of ANNs

- McCulloch and Pitts
- Donald Hebb “The Organization of Behaviour”
- ADALINE for binary patterns by Widrow and Hoff
- Rosenblatt’s perceptron






After Rosenblatt (1958)

- A historical note
- Course outline
- Introduction to ANNs

Brief history of ANNs


- McCulloch and Pitts
- Donald Hebb “The Organization of Behaviour”
- ADALINE for binary patterns by Widrow and Hoff
- Rosenblatt’s perceptron
- Marvin Minsky and Seymour Papert's criticism (1969)
- first AI winter: fears and outrageous claims led to lower interest and poor funding



- **A historical note**
- Course outline
- Introduction to ANNs

Brief history of ANNs

- Hopfield's impact



- **A historical note**
- Course outline
- Introduction to ANNs


Brief history of ANNs

- Hopfield's impact
- Renewed enthusiasm in the 1980s and 1990s (backpropagation, 1986)

- **A historical note**
- Course outline
- Introduction to ANNs

Brief history of ANNs

- Hopfield's impact
- Renewed enthusiasm in the 1980s and 1990s (backpropagation, 1986)
- interest shifts towards more mathematically rigorous statistical learning theory
- another revival of interest in deep neural networks



A wide scope of applications

- Pattern recognition: classification and clustering
- General interpolation problems (generalisation)
- Data representations, coding and compression
- Signal processing
- Time series prediction
- System identification
- Decision support (e.g. medical or industrial diagnostics)
- Memory storage, modelling
- Optimisation, combinatorial problems
- etc.....

Course outline

- Learning activities
 - 11 lecture themes + 1 summary and Q&A lecture
 - 10 “lab” sessions (evaluation and support)
 - 4 lab assignments (in groups of 3 students)

Course outline

- Learning activities
 - 11 lecture themes + 1 summary and Q&A lecture
 - 10 “lab” sessions (evaluation and support)
 - 4 lab assignments (in groups of 3 students)
- Assessment, grading
 - 4 mandatory lab assignments with bonus point deadlines (check in Canvas: under Lab Review as well as Content, assessment)
 - lab review (redovisning) and a short report to submit in Canvas (P/F)
 - written exam: A-F

Lab review

- be prepared for short 10-min demonstration of the lab results (in groups) and 2-3 min questions, we do not tend to inspect your code
- time is short so planning and preparation are critical –
 - focus on main points, key questions – what is the lab about?
 - what are the assumptions?
 - share your insights and formulate clear conclusions
 - what did you learn?
 - are there any open questions left?
- you are in charge of the lab review but we are likely to ask questions
- please ensure the entire group is involved (time sharing)
- please demonstrate in the form of a report with clear & clean figures

Lab report

- a concise document to be uploaded to Canvas as a pdf file before your lab demonstration (bring the report with you)
- a typical layout involves (for each lab there is a template)
 - clear header (lab title, course signature, list of authors)
 - aim/scope of the lab assignment, assumptions, tools used
 - main results/findings supported with figures or tables (with captions, and always referred to in the main text) along with short commentary, your interpretations
 - reflections, open questions and conclusions (to the point)
 - no code needed and please do not add any appendices – there is an upper limit for the number of pages, e.g. 6 for lab 1.
- follow instructions on Canvas

Course content

Lectures

- L1: Course introduction and fundamental concepts
- L2: From perceptron learning rule to backpropagation in feedforward networks – supervised learning
- L3: Generalization, regularization, model selection and validation
- L4: Practical aspects of ANN approaches to pattern recognition problems
- L5: Radial basis function networks and introduction to unsupervised learning
- L6: Self-organising maps


Course content

Lectures

- L7: Temporal processing with ANNs: feedforward vs recurrent network architectures
- L8: Hopfield networks and introduction to stochastic networks
- L9: Deep learning fundamentals: general philosophy and a review of deep architectures
- L10: Representation learning and deep generative models
- L11: Deep neural networks: practicalities, challenges and current trends
- L12: Course summary, old exam questions, Q&A



Fundamental characteristics

- **nodes, units**
- activation function
- learning rule
- topology, network architecture
- data





Fundamental characteristics

- nodes, units
- **activation function**
- learning rule
- topology, network architecture
- data




Fundamental characteristics

- nodes, units
- **activation function**
- learning rule
- topology, network architecture
- data



Fundamental characteristics


- nodes, units
- activation function
- **learning rule**
- topology, network architecture
- data



Fundamental characteristics


- nodes, units
- activation function
- **learning rule**
- topology, network architecture
- data

unsupervised,
competitive
out






Fundamental characteristics

- nodes, units
- activation function
- **learning rule**
- topology, network architecture
- data




Fundamental characteristics

- nodes, units
- activation function
- learning rule
- **topology, network architecture**
- data



Fundamental characteristics


- nodes, units
- activation function
- learning rule
- topology, network architecture
- data



- A historical note
- Course outline
- **Introduction to ANNs**


Biomimetic nature of ANNs

Inspirations from biology



electrically excitable cells for
processing and transmitting
information
(ca. 100 billion in the human brain)


Multitude of neuron types in the brain




- A historical note
- Course outline
- **Introduction to ANNs**

Biomimetic nature of ANNs

Inspirations from biology





electrically excitable cells for processing and transmitting information
(ca. 100 billion in the human brain)



- A historical note
- Course outline
- **Introduction to ANNs**


Biomimetic nature of ANNs

Inspirations from biology – action potential (firing)



Biomimetic nature of ANNs


Inspirations from biology – action potential (firing)



- commonly referred to as “spike” (nerve impulse)
- threshold phenomenon accounts for “all-or-nothing” paradigm
- spike (action potential of fixed amplitude) travels along the axon

Biomimetic nature of ANNs

Inspirations from biology – action potential (firing)





- commonly referred to as “spike” (nerve impulse)
- threshold phenomenon accounts for “all-or-nothing” paradigm
- spike (action potential of fixed amplitude) travels along the axon
- when spike reaches synaptic terminal, it contributes to post-synaptic potentials


- A historical note
- Course outline
- **Introduction to ANNs**

Biomimetic nature of ANNs

Inspirations from biology – action potential (firing)



Depolarization
(Excitatory)



Hyperpolarization
(Inhibitory)

Biomimetic nature of ANNs

Biological neurons and networks – computational aspects

- numerous simple units (type of integrators)
- binary communication – spike (“all-or-nothing”)
- but analog transmission in synapse (connection weight)
- many inputs, interconnect
- summation (integration) of inputs

Biomimetic nature of ANNs

Biological neurons and networks – computational aspects

- local processing, learning (based on local information)
- synaptic memory (memory stored in weights)
- synaptic plasticity – weight adaptation according to learning rules (largely phenomenological)
- weights can be positive and negative (but Dale's law)
- fast nature of parallel processing
- tolerance to errors/noise in data, weights etc.