



DD2437 – Artificial Neural Networks and Deep Architectures (annda)

Lecture 8b: **Boltzmann machines and RBMs, autoencoders**

Pawel Herman

Computational Science and Technology (CST)
KTH Royal Institute of Technology

September 2019


- Associative memory
- Hopfield networks
- Memory storage and TSP example
- Stochastic networks – Boltzmann machine

Lecture overview

- Boltzmann machine
- Restricted Boltzmann machine, RBM
- Autoencoders

- Associative memory
- **Hopfield networks**
- Memory storage and TSP example
- Stochastic networks – Boltzmann machine

Hopfield network



$$\forall_i w_{i,i} = 0 \quad \text{no self-connections}$$

$$\vec{x}' = \text{sgn}(\mathbf{W}\vec{x} + \vec{\theta})$$

$$E(\text{state} = \vec{x}) = -\frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n w_{i,j} x_i x_j + \sum_{i=1}^n \theta_i x_i$$

- Associative memory
- Hopfield networks
- Memory storage and TSP example
- **Stochastic networks – Boltzmann machine**

From Hopfield networks to Boltzmann machines

- Continuous Hopfield network

$$x_i = \frac{1}{1 + e^{-a}} \quad \text{instead of} \quad x_i = \text{sgn}(a)$$

- Stochastic component

$$x_i = \begin{cases} 1 & \text{with probability } p_i \\ -1, & \text{with probability } 1-p_i \end{cases} \quad p_i = \frac{1}{1 + e^{-\frac{1}{T} \sum_j w_{i,j} x_j}}$$

T is a positive temperature const.

- Associative memory
- Hopfield networks
- Memory storage and TSP example
- **Stochastic networks – Boltzmann machine**

From Hopfield networks to Boltzmann machines

- Stochastic component

$$x_i = \begin{cases} 1 & \text{with probability } p_i \\ -1, & \text{with probability } 1-p_i \end{cases}$$

$$p_i = \frac{1}{1 + e^{-\frac{1}{T} \sum_j w_{i,j} x_j}}$$

$$p(\nu) = \frac{1}{1 + e^{-\nu}} \quad \text{where} \quad \nu = \frac{1}{T} \sum_j w_{i,j} x_j$$

T controls the level of randomness

- Associative memory
- Hopfield networks
- Memory storage and TSP example
- **Stochastic networks – Boltzmann machine**



From Hopfield networks to Boltzmann machines

- Stochastic component

$$x_i = \begin{cases} 1 & \text{with probability } p_i \\ -1, & \text{with probability } 1-p_i \end{cases}$$

$$p_i = \frac{1}{1 + e^{-\frac{1}{T} \sum_j w_{i,j} x_j}}$$

$$p(\nu) = \frac{1}{1 + e^{-\nu}} \quad \text{where} \quad \nu = \frac{1}{T} \sum_j w_{i,j} x_j$$



- Associative memory
- Hopfield networks
- Memory storage and TSP example
- **Stochastic networks – Boltzmann machine**


From Hopfield networks to Boltzmann machines

- Stochastic component

$$x_i = \begin{cases} 1 & \text{with probability } p_i \\ -1, & \text{with probability } 1-p_i \end{cases}$$

$$p_i = \frac{1}{1 + e^{-\frac{1}{T} \sum_j w_{i,j} x_j}}$$

Analogy to **simulated annealing** (relaxation technique common in metallurgy)




$$p_{\Delta E} = \frac{1}{1 + e^{\Delta E/T}}$$

“When optimising a large complex system with many degrees of freedom, instead of always going downhill, try to go downhill most of the time”

- Associative memory
- Hopfield networks
- Memory storage and TSP example
- **Stochastic networks – Boltzmann machine**

Simulated annealing to reach the global energy min



The critical role of temperature T .



- Associative memory
- Hopfield networks
- Memory storage and TSP example
- **Stochastic networks – Boltzmann machine**

Simulated annealing to reach the global energy min

The critical role of temperature T .



$$p_{\Delta E} = \frac{1}{1 + e^{\Delta E / T}}$$

- Associative memory
- Hopfield networks
- Memory storage and TSP example
- **Stochastic networks – Boltzmann machine**

From Hopfield networks to Boltzmann machines

- Energy of this stochastic network is the same as before

$$E = -\frac{1}{2} \vec{x}^T \mathbf{W} \vec{x} = -\frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n w_{i,j} x_i x_j$$

- The key difference is a stochastic nature of transitions

from state $s1$ to $s2$:

$$p_{s1 \rightarrow s2} = \frac{1}{1 + e^{(E_2 - E_1)/T}} = \frac{1}{1 + e^{\Delta E/T}}$$

- Associative memory
- Hopfield networks
- Memory storage and TSP example
- **Stochastic networks – Boltzmann machine**

From Hopfield networks to Boltzmann machines

- Energy of this stochastic network is the same as before


$$E = -\frac{1}{2} \vec{x}^T \mathbf{W} \vec{x} = -\frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n w_{i,j} x_i x_j$$

- Given a set of examples $\{\vec{x}_i\}_1^m$ the idea is to adjust \mathbf{W} to describe data distribution (well matched to these examples)

$$P(\vec{x} | \mathbf{W}) = \frac{e^{-E}}{Z} = \frac{1}{Z(\mathbf{W})} \exp\left(-\frac{1}{2} \vec{x}^T \mathbf{W} \vec{x}\right)$$

- Associative memory
- Hopfield networks
- Memory storage and TSP example
- **Stochastic networks – Boltzmann machine**


Hidden and visible units



- Symmetric connections between visible, v , and hidden neurons, h
- Hidden neurons help account for higher-order correlations in the input vectors (data)
- Visible units provide interface to the external world – environment (data, $v=x$)
- Hidden units operate freely and are used to explain environmental input vectors

- Associative memory
- Hopfield networks
- Memory storage and TSP example
- **Stochastic networks – Boltzmann machine**


Hidden and visible units



- Symmetric connections between visible, v , and hidden neurons, h
- Hidden neurons help account for higher-order correlations in the input vectors (data)
- Visible units provide interface to the external world – environment (data, $v=x$)
- Hidden units operate freely and are used to explain environmental input vectors

- Associative memory
- Hopfield networks
- Memory storage and TSP example
- **Stochastic networks – Boltzmann machine**

Hidden and visible units



$$v^{(p)} = x^{(p)}$$

$$\Downarrow$$

$$y^{(p)} = [x^{(p)}, h]$$

- Symmetric connections between visible, v , and hidden neurons, h
- Hidden neurons help account for higher-order correlations in the input vectors (data)
- Visible units provide interface to the external world – environment (data, $v=x$)
- Hidden units operate freely and are used to explain environmental input vectors
- Modelling a probability distribution (and hidden representation) by clamping patterns onto the visible units $v^{(p_i)} = x^{(p_i)}$

- Associative memory
- Hopfield networks
- Memory storage and TSP example
- **Stochastic networks – Boltzmann machine**

Boltzmann learning

- The primary goal is to correctly model input patterns according to Boltzmann distribution
 - each input pattern is assumed to last long enough (it might have to be clamped for long) for the network to reach thermal equilibrium (converge) at temperature T
 - to reduce this time, simulated annealing is used with a sequence decreasing temperatures (from “hot” to “cold”)
- Essentially, hidden units learn probabilistically representation of data (seen through visible units)

- Associative memory
- Hopfield networks
- Memory storage and TSP example
- **Stochastic networks – Boltzmann machine**

Boltzmann learning

- The idea is to maximise log-likelihood, $L(\mathbf{W}) = \log (P(\mathbf{X})|\mathbf{W})$

$$\Delta w_{ji} = \varepsilon \frac{\partial L(\mathbf{W})}{\partial w_{ji}} = \eta (\rho_{j,i}^+ - \rho_{j,i}^-), \quad \eta = \varepsilon / T$$

- Associative memory
- Hopfield networks
- Memory storage and TSP example
- **Stochastic networks – Boltzmann machine**

Boltzmann learning

- The idea is to maximise log-likelihood, $L(\mathbf{W}) = \log(P(\mathbf{X})|\mathbf{W})$

$$\Delta w_{ji} = \varepsilon \frac{\partial L(\mathbf{W})}{\partial w_{ji}} = \eta(\rho_{j,i}^+ - \rho_{j,i}^-), \quad \eta = \varepsilon / T$$

$$\frac{\partial L(\mathbf{W})}{\partial w_{i,j}} = \frac{\partial}{\partial w_{i,j}} \log P(\{\mathbf{x}^{(p)}\}_1^M | \mathbf{W}) = \sum_p \left\{ \underbrace{\langle y_i y_j \rangle_{P(\mathbf{h}|\mathbf{v}=\mathbf{x}^{(p)}, \mathbf{W})}}_{\text{positive phase (awake), with clamping, } \mathbf{v}^{(p)}=\mathbf{x}^{(p)}} - \underbrace{\langle y_i y_j \rangle_{P(\mathbf{v}, \mathbf{h}|\mathbf{W})}}_{\text{negative phase (sleep) free running}} \right\}$$

$$\langle y_i y_j \rangle_{P(\mathbf{h}|\mathbf{v}=\mathbf{x}^{(p)}, \mathbf{W})} = \sum_p \sum_h P(\mathbf{h} | \mathbf{v} = \mathbf{x}^{(p)}) y_i y_j$$

$$\langle y_i y_j \rangle_{P(\mathbf{v}, \mathbf{h}|\mathbf{W})} = \sum_p \sum_y P(\mathbf{y}) y_i y_j$$

- Associative memory
- Hopfield networks
- Memory storage and TSP example
- **Stochastic networks – Boltzmann machine**

Boltzmann learning – two phases

$$\frac{\partial L(\mathbf{W})}{\partial w_{i,j}} = \frac{\partial}{\partial w_{i,j}} \log P(\{\mathbf{x}^{(p)}\}_1^M \mid \mathbf{W}) = \sum_p \left\{ \langle y_i y_j \rangle_{P(\mathbf{h}|\mathbf{v}=\mathbf{x}^{(p)}, \mathbf{W})} - \langle y_i y_j \rangle_{P(\mathbf{v}, \mathbf{h}|\mathbf{W})} \right\}$$

$$\Delta w_{i,j} \propto \langle y_i, y_j \rangle_{\text{data}} - \langle y_i, y_j \rangle_{\text{model}}$$

- Associative memory
- Hopfield networks
- Memory storage and TSP example
- **Stochastic networks – Boltzmann machine**

Boltzmann learning – two phases

$$\frac{\partial L(\mathbf{W})}{\partial w_{i,j}} = \frac{\partial}{\partial w_{i,j}} \log P(\{\mathbf{x}^{(p)}\}_1^M | \mathbf{W}) = \sum_p \left\{ \langle y_i y_j \rangle_{P(\mathbf{h}|\mathbf{v}=\mathbf{x}^{(p)}, \mathbf{W})} - \langle y_i y_j \rangle_{P(\mathbf{v}, \mathbf{h}|\mathbf{W})} \right\}$$

$$\Delta w_{i,j} \propto \langle y_i, y_j \rangle_{\text{data}} - \langle y_i, y_j \rangle_{\text{model}}$$

- **Positive** phase implies clamping the inputs (relative fast)

$$\langle y_i, y_j \rangle_{\text{data}} \longleftarrow \text{Expected value at thermal equilibrium}$$

- **Negative** phase involves updating all the units (can be very slow)

$$\langle y_i, y_j \rangle_{\text{model}}$$

- Associative memory
- Hopfield networks
- Memory storage and TSP example
- **Stochastic networks – Boltzmann machine**

Boltzmann learning – two phases

$$\frac{\partial L(\mathbf{W})}{\partial w_{i,j}} = \frac{\partial}{\partial w_{i,j}} \log P(\{\mathbf{x}^{(p)}\}_1^M | \mathbf{W}) = \sum_p \left\{ \langle y_i y_j \rangle_{P(\mathbf{h}|\mathbf{v}=\mathbf{x}^{(p)}, \mathbf{W})} - \langle y_i y_j \rangle_{P(\mathbf{v}, \mathbf{h}|\mathbf{W})} \right\}$$

$$\Delta w_{i,j} \propto \langle y_i, y_j \rangle_{\text{data}} - \langle y_i, y_j \rangle_{\text{model}}$$

- **Positive** phase implies clamping the inputs (relative fast)

Thermal equilibrium does not imply only that the system settles down into the lowest energy state.

Expected value at thermal equilibrium

- **Nega**

It is about the convergence of probability distribution over different configurations.

- Associative memory
- Hopfield networks
- Memory storage and TSP example
- **Stochastic networks – Boltzmann machine**

Boltzmann learning – two phases

$$\frac{\partial L(\mathbf{W})}{\partial w_{i,j}} = \frac{\partial}{\partial w_{i,j}} \log P(\{\mathbf{x}^{(p)}\}_1^M | \mathbf{W}) = \sum_p \left\{ \langle y_i y_j \rangle_{P(\mathbf{h}|\mathbf{v}=\mathbf{x}^{(p)}, \mathbf{W})} - \langle y_i y_j \rangle_{P(\mathbf{v}, \mathbf{h}|\mathbf{W})} \right\}$$

$$\Delta w_{i,j} \propto \langle y_i, y_j \rangle_{\text{data}} - \langle y_i, y_j \rangle_{\text{model}}$$

- **Positive** phase implies clamping the inputs (relative fast)

“Hebbian learning” $\langle y_i, y_j \rangle_{\text{data}}$


- **Negative** phase involves updating all the units (can be very slow)

“Hebbian forgetting” $\langle y_i, y_j \rangle_{\text{model}}$ prevent from learning false, spontaneously generated states


- Associative memory
- Hopfield networks
- Memory storage and TSP example
- **Stochastic networks – Boltzmann machine**

Next step to decrease the complexity

- Restricted Boltzmann machines




- Deep belief networks




- Data representations
- **Restricted Boltzmann machine**
- Autoencoders
- Deep generative models

Computational blocks for learning representations

- Two key approaches to greedy layer-wise pretraining
 - regularized Boltzmann machines (RBMs)

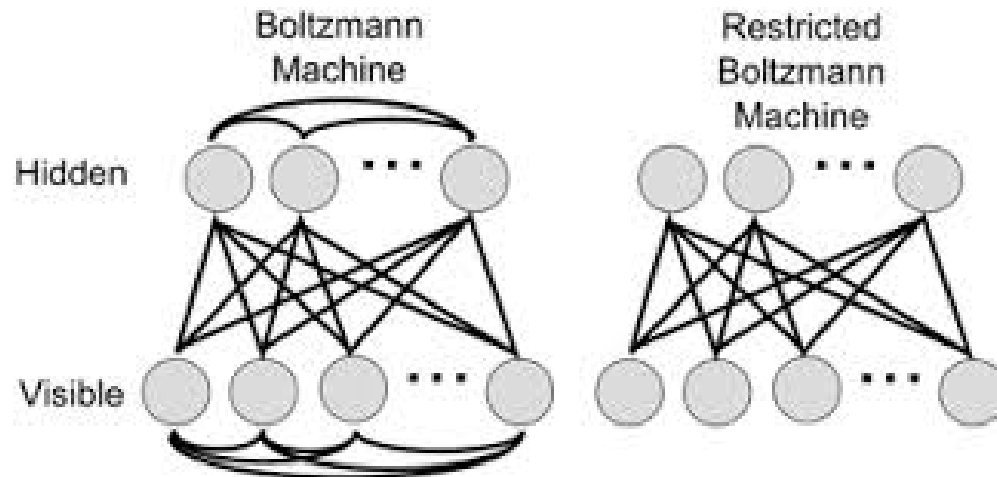


- autoencoders



- Data representations
- **Restricted Boltzmann machine**
- Autoencoders
- Deep generative models

Restricted Boltzmann machine (RBM)




Visible and hidden units are conditionally independent given one another

$$p(\mathbf{h} | \mathbf{v}) = \prod_i p(h_i | \mathbf{v})$$

$$p(\mathbf{v} | \mathbf{h}) = \prod_j p(v_j | \mathbf{h})$$

- Data representations
- **Restricted Boltzmann machine**
- Autoencoders
- Deep generative models

Restricted Boltzmann machine (RBM)



Visible and hidden units are conditionally independent given one another

$$p(\mathbf{h} | \mathbf{v}) = \prod_i p(h_i | \mathbf{v})$$

$$p(\mathbf{v} | \mathbf{h}) = \prod_j p(v_j | \mathbf{h})$$

Following the same principle of maximising log likelihood by means of gradient ascent, one obtains:

$$\Delta w_{ji} = \varepsilon \frac{\partial L(\mathbf{W})}{\partial w_{ji}} = \varepsilon \left(\langle v_j h_i \rangle_{\text{data}} - \langle v_j h_i \rangle_{\text{model}} \right)$$


- Data representations
- **Restricted Boltzmann machine**
- Autoencoders
- Deep generative models

Restricted Boltzmann machine (RBM)

Visible and hidden units are conditionally independent given one another

$$P(h_i = 1 | \mathbf{v}) = \frac{1}{1 + \exp(-bias_{h_i} - \mathbf{v}^T \mathbf{W}_{:,i})}$$

$$P(v_j = 1 | \mathbf{h}) = \frac{1}{1 + \exp(-bias_{v_j} - \mathbf{W}_{j,:} \mathbf{h})}$$



$$p(\mathbf{h} | \mathbf{v}) = \prod_i p(h_i | \mathbf{v})$$

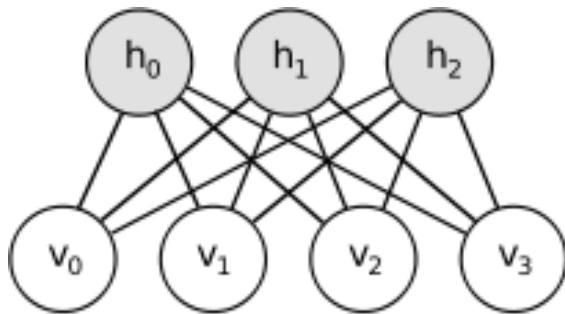
$$p(\mathbf{v} | \mathbf{h}) = \prod_j p(v_j | \mathbf{h})$$

Following the same principle of maximising log likelihood by means of gradient ascent, one obtains:

$$\Delta w_{ji} = \varepsilon \frac{\partial L(\mathbf{W})}{\partial w_{ji}} = \varepsilon \left(\langle v_j h_i \rangle_{\text{data}} - \langle v_j h_i \rangle_{\text{model}} \right)$$

- Data representations
- **Restricted Boltzmann machine**
- Autoencoders
- Deep generative models




RBM learning with Contrastive Divergence (CD)



$$P(h_i = 1 | \mathbf{v}) = \frac{1}{1 + \exp(-bias_{h_i} - \mathbf{v}^T \mathbf{W}_{:,i})}$$

$$P(v_j = 1 | \mathbf{h}) = \frac{1}{1 + \exp(-bias_{v_j} - \mathbf{W}_{j,:} \mathbf{h})}$$


Gibbs sampling



Hinton, 2003

- Data representations
- **Restricted Boltzmann machine**
- Autoencoders
- Deep generative models

RBM learning with Contrastive Divergence (CD)






$$P(h_i = 1 | \mathbf{v}) = \frac{1}{1 + \exp(-bias_{h_i} - \mathbf{v}^T \mathbf{W}_{:,i})}$$

$$P(v_j = 1 | \mathbf{h}) = \frac{1}{1 + \exp(-bias_{v_j} - \mathbf{W}_{j,:} \mathbf{h})}$$

GOOD TO KNOW:

Contrastive Divergence does not optimise the likelihood but it works effectively!

Gibbs sampling




Hinton, 2003

- Data representations
- **Restricted Boltzmann machine**
- Autoencoders
- Deep generative models

CD_k recipe for training RBM

Gibbs sampling



- 1) Clamp the visible units with an input vector and update hidden units.

$$P(h_i = 1 | \mathbf{v}) = \left(1 + \exp \left(-bias_{h_i} - \mathbf{v}^T \mathbf{W}_{:,i} \right) \right)^{-1}$$

- 2) Update all the visible units in parallel to get a **reconstruction**.

$$P(v_j = 1 | \mathbf{h}) = \left(1 + \exp \left(-bias_{v_j} - \mathbf{W}_{j,:} \mathbf{h} \right) \right)^{-1}$$

- 3) Collect the statistics for correlations after k steps using mini-batches and update weights:

$$\Delta w_{j,i} = \frac{1}{N} \sum_{n=1}^N \left(v_j^{(n)} h_i^{(n)} - \hat{v}_j^{(n)} \hat{h}_i^{(n)} \right)$$

- Data representations
- **Restricted Boltzmann machine**
- Autoencoders
- Deep generative models

From RBM to Gaussian-Bernoulli RBM

Bernoulli-Bernoulli (binary-binary)

$$p(v_i = 1|\mathbf{h}) = g\left(\sum_j W_{ij}b_j + b_i\right)$$

$$p(b_j = 1|\mathbf{v}) = g\left(\sum_i W_{ij}v_i + a_j\right)$$



Gaussian-Bernoulli (real/cont.-binary)

$$p(v_i = x|\mathbf{h}) = \frac{1}{\sqrt{2\pi}\sigma_i} \exp\left(-\frac{\left(x - b_i - \sigma_i \sum_j b_j W_{ij}\right)^2}{2\sigma_i^2}\right),$$

$$p(b_j = 1|\mathbf{v}) = g\left(b_j + \sum_i W_{ij} \frac{v_i}{\sigma_i}\right),$$



Visible units are real-valued whereas hidden units remain binary.

Salakhutdinov, 2015

- Data representations
- **Restricted Boltzmann machine**
- Autoencoders
- Deep generative models

From RBM to Gaussian-Bernoulli RBM

Bernoulli-Bernoulli (binary-binary)

$$p(v_i = 1|\mathbf{h}) = g\left(\sum_j W_{ij}b_j + b_i\right)$$

$$p(b_j = 1|\mathbf{v}) = g\left(\sum_i W_{ij}v_i + a_j\right)$$



Gaussian-Bernoulli (real/cont.-binary)

$$p(v_i = x|\mathbf{h}) = \frac{1}{\sqrt{2\pi}\sigma_i} \exp\left(-\frac{\left(x - b_i - \sigma_i \sum_j b_j W_{ij}\right)^2}{2\sigma_i^2}\right),$$

$$p(b_j = 1|\mathbf{v}) = g\left(b_j + \sum_i W_{ij} \frac{v_i}{\sigma_i}\right),$$

Visible units are real-valued whereas hidden units remain binary.


The derivative of the log-likelihood:

$$\frac{\partial \log P(\mathbf{v}; \theta)}{\partial W_{ij}} = \mathbb{E}_{P_{\text{data}}} \left[\frac{1}{\sigma_i} v_i b_j \right] - \mathbb{E}_{P_{\text{model}}} \left[\frac{1}{\sigma_i} v_i b_j \right]$$

Salakhutdinov, 2015

- Data representations
- Restricted Boltzmann machine
- **Autoencoders**
- Deep generative models

Autoencoders – principles



Decoder: $x' = f(x)$


Encoder: $h = g(f(x))$

The idea is to minimise the loss function, L :

$$L(x, g(f(x)))$$

- Data representations
- Restricted Boltzmann machine
- **Autoencoders**
- Deep generative models

Autoencoders



Encourage h to give small reconstruction error:

- e.g. $Loss = \sum_m \|x^{(m)} - DECODER(ENCODER(x^{(m)}))\|^2$
- Reconstruction: $x' = \sigma(W'\sigma(Wx + b) + d)$

Different types of autoencoders

- Undercomplete autoencoders
 - hidden layer is smaller than the input dimensionality
- Overcomplete regularised autoencoders
 - Larger hidden layer size with the regularisation (to avoid overfitting and copying input to the output)

$$L(x, g(f(x))) + \Omega(h), \quad h = f(x)$$

- Sparse autoencoders, denoising autoencoders
- Deep autoencoders

Sparse autoencoders

- Penalizing non-sparse solutions can be seen as adding latent variables with a prior and maximising likelihood


$$\log p_{\text{model}}(\mathbf{h}, \mathbf{x}) = \log p_{\text{model}}(\mathbf{h}) + \log p_{\text{model}}(\mathbf{x} | \mathbf{h})$$

for example: $p_{\text{model}}(\mathbf{h}) = \prod_i \frac{\lambda}{2} e^{-\lambda |h_i|} \Rightarrow \Omega(\mathbf{h}) = \lambda \sum |h_i|$

- Data representations
- Restricted Boltzmann machine
- **Autoencoders**
- Deep generative models

Denoising autoencoders

$$L(\mathbf{x}, g(f(\tilde{\mathbf{x}}))), \quad h = f(\tilde{\mathbf{x}})$$




Corrupted copy of x

Autoencoders have to undo this corruption beyond simply coping the input.

Denoising autoencoders

$$L(\mathbf{x}, g(f(\tilde{\mathbf{x}}))), \quad h = f(\tilde{\mathbf{x}})$$



Corrupted copy of x

Autoencoders have to undo this corruption beyond simply coping the input.

1. A training sample is sampled from the training data.
2. A corrupted version of the sample \mathbf{x} is drawn from some corruption process

$$C(\tilde{\mathbf{x}} \mid \mathbf{x} = \mathbf{s})$$

3. $(\mathbf{x}, \tilde{\mathbf{x}})$ is used as a training sample to estimate the autoencoder's reconstruction distribution $p_{reconstruction}(\tilde{\mathbf{x}} \mid \mathbf{x}) = p_{decoder}(\mathbf{x} \mid \mathbf{h}), \quad \mathbf{h} = f(\tilde{\mathbf{x}})$

- Data representations
- Restricted Boltzmann machine
- **Autoencoders**
- Deep generative models


Denoising autoencoders

$$L(\mathbf{x}, g(f(\tilde{\mathbf{x}}))), \quad h = f(\tilde{\mathbf{x}})$$

Corrupted copy of x

Autoencoders have to undo this corruption beyond simply coping the input.

Learning a *vector field* around a *low-dimensional manifold* . . .



. . . with the principle that only the variations tangent to the manifold around \mathbf{x} should be accounted for by changes in \mathbf{h}



Goodfellow et al.

- Data representations
- Restricted Boltzmann machine
- **Autoencoders**
- Deep generative models

Denoising autoencoders

When training autoencoders there is a **compromise**

- I. Need to approximately recover \mathbf{x} – *reconstruction* force
- II. Need to satisfy the regularization term – *regularisation* force.



Goodfellow et al.